



Offline RL on logged agent traces: reward-label real runs → evaluate fixes → warm-start & explore, with zero new API calls

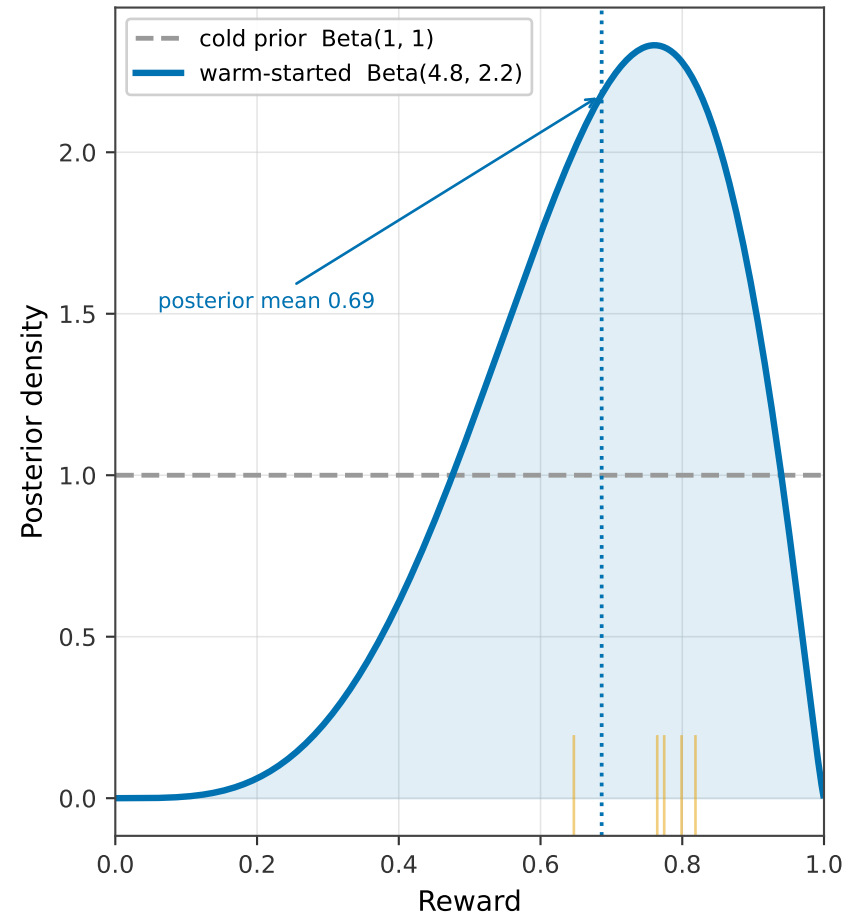
A · Offline policy evaluation of the harness fixes



B · Trajectory-rooted best-of-N (zero new calls)



C · Warm-start the N=1 arm from logged rewards



Trajectories harvested once from our own LangSmith logs (no model inference); all analysis runs offline against the committed cache using the physically-gated strap.eval reward and strap.learn best-of-N / Thompson bandit.